

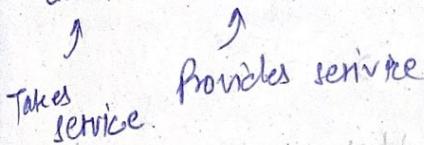
Networks :

Make M/Cs communicate

TCP, UDP

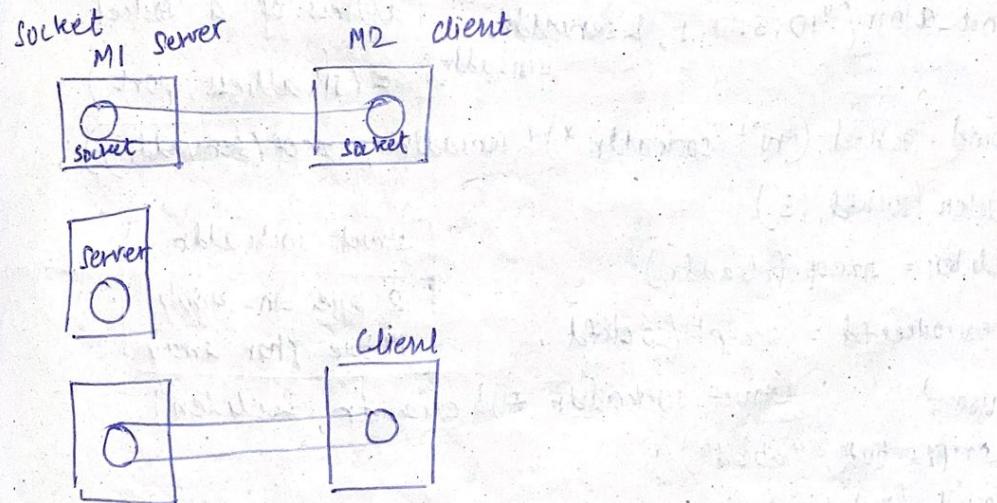
TCP - FIFO, reliable communication b/w "2-endpoints"
bit byte-oriented

Client - Server



Server :

TCP comm.



Server:

socket() — opens a socket

bind() — gives an address to the socket

listen()

accept() — waits for a client to connect

send() / receive() — once a client connects

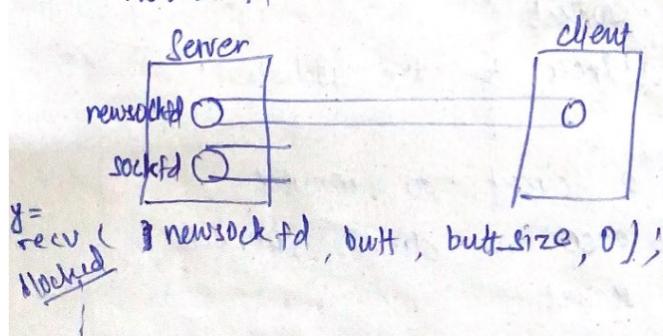
close() — closes the socket

Client :

socket() - opens a socket
bind() ?? not needed ??
connect() - connect to server
send() / recv() data
close()

```
int sockfd;
sockfd = socket(family, type, protocol)
    /   |
    AF_INET  SOCK_STREAM  0
struct sockaddr_in serveraddr, cliaddr;
serveraddr.sin_family = AF_INET;
serveraddr.sin_port = htons(20000); // network byte order
serveraddr.sin_addr
inet_aton("10.5.17.1", &serveraddr.sin_addr) // address of a socket
                                                = (IP address, port)
bind(sockfd, (struct sockaddr *) &serveraddr, sizeof(serveraddr));
listen(sockfd, 5)
clilen = sizeof(cliaddr)
newsocketfd = accept(sockfd,
    blocked          (struct sockaddr * ) &cliaddr, &clilen)
strcpy(buff, "abcd");
send(newsocketfd, buff, strlen(buff)+1, 0);
newsocketfd
```

struct sockaddr
2 byte .in-family
14 byte char array



y = recv(newsocketfd, buff, buff_size, 0);

blocked

Networks

UDP

correctedness, unreliable, no FIFO

Message-oriented



① What you send in one `send()` will be received in ~~one~~ ~~one~~ one `recv()` if received.

② If the buffer size in `recv()` call < what is sent, remaining data will be lost

UDP server

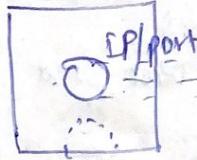
`socket(_____, SOCK_DGRAM, __)`

`bind()`

`sendto() / recvfrom`

`close()`

Server



UDP client

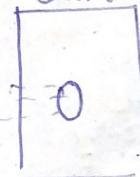
`socket(_____, SOCK_DGRAM, __)`

:

`send() / recv` \equiv `sendto() / recvfrom`

`close()`

Client

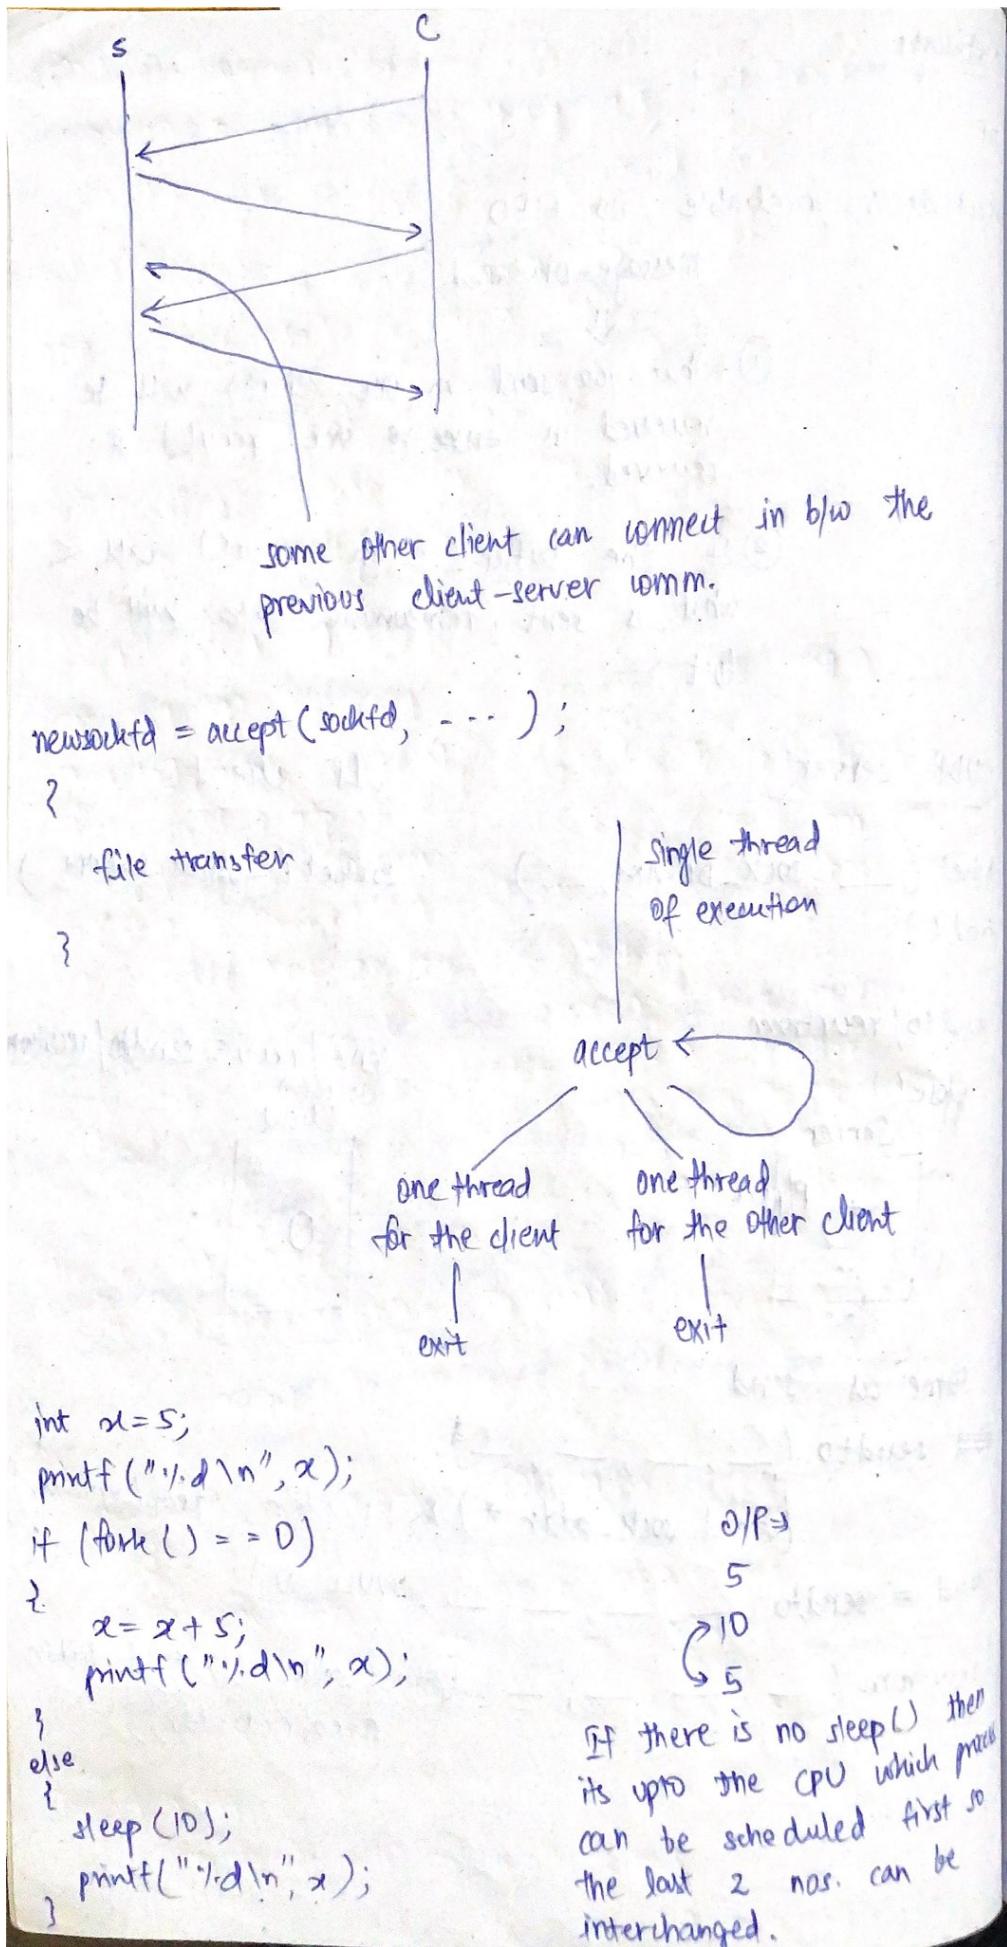


Same as `send`

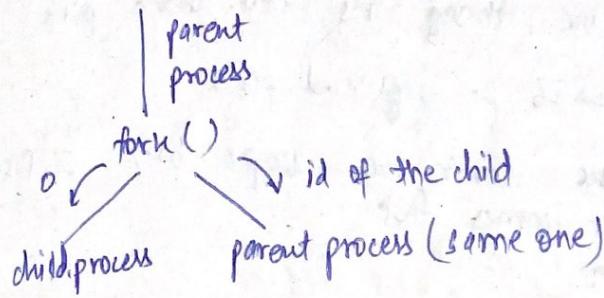
~~sendto~~.`(_____, _____, _____, _____,`
~~(struct sockaddr *)~~ `&servaddr, sizeof(servaddr);`

`send = sendto(_____, _____, _____, _____, NULL, 0);`

`recvfrom(_____, _____, _____, _____, (struct sockaddr *) &cliaddr,`
`sizeof(cliaddr));`



$x = \text{fork}()$



The memory gets copied in both the processes.

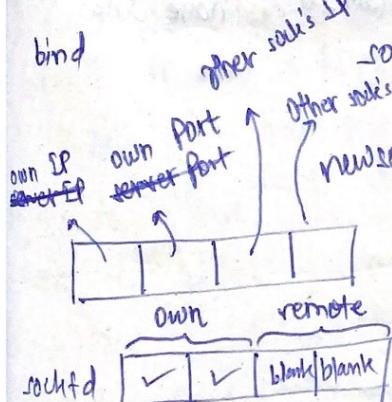


```

while(1) {
    newsockfd = accept (sockfd, —, —, —);
    if (fork () == 0) {
        close (sockfd);
        comm. with client
        exit (0);
    }
    close (newsockfd);
}
  
```

socket (-)

bind



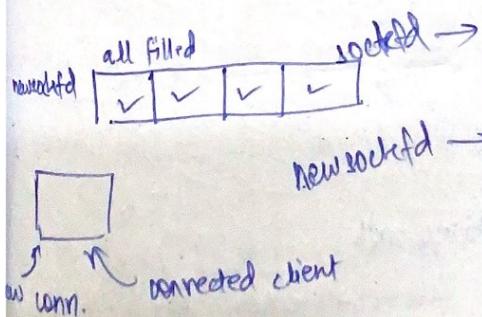
sockfd	✓	✓	blank	blank
--------	---	---	-------	-------

Per process Table

Id

System wide Table

Info socket



X close()

Ports

0 - 1023 — cannot use these ports (well known ports)

1024 - $\frac{3}{4}$ th of 65536 — registered ports

If we want to make server to accept both TCP & UDP client whichever comes first

We have a function for this
poll / select

```

struct pollfd fdset[2];
fdset[0].fd = tcp.sockfd;
fdset[0].events = POLLIN;
fdset[1].fd = udp.sockfd;
fdset[1].events = POLLIN;
poll(fdset, 2, timeout);
ret = poll(fdset, 2, timeout); //blocking

```

pollfd

{ int fd;

short events;

short revents;

}

ret < 0 \Rightarrow error

ret = 0 \Rightarrow timeout

ret > 0 \Rightarrow x of the socket
(say x) have data

else if (ret > 0) {

if (fdset[0].revents == POLLIN)

{ accept(); }

if (fdset[1].revents == POLLIN)

{

recvfrom(); }

}

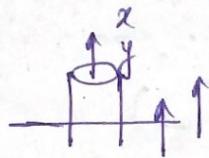
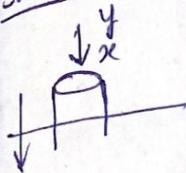
}

timeout \Rightarrow int

Signifies timeout in milliseconds

Networks:

socket



TCP socket

Reliable

In order delivery

Byte oriented

UDP socket

Unreliable

May be out of order

Message oriented

TCP server

waits for client

TCP Client

asks for service

Client - server programming

TCP Server

Create socket - `socket()`

Bind the socket - `bind()`
(specify IP address & port)

One m/c has one IP addr only, difference is in port

Configure queue - `listen()`

Wait for client - `accept()`

Recv for client - `recv()`

Send to client - `send()`

TCP Client

Create socket - `socket()`

No bind -

BUT every socket has an IP, port
Because no one externally connects client IP, port is

assigned dynamically

connect() → accept() from server side

socket (__ , __ , __)
↓ ↓ ↓
AF-INET SOCK-STREAM
OR (for TCP socket)
PF-INET

(using the internet family
of sockets)

socketfd = socket(...)

↑
Returns socket descriptor

err = bind (socketfd, (struct sockaddr *) servaddr, sizeof
(struct sockaddr));

struct sockaddr_in serv_addr;

serv_addr.sin_family = AF-INET

serv_addr.sin_port = 20000; (take anything > 10000)

serv_addr.sin_addr.s_addr = INADDR-ANY

(accept packets from any IP
address)

Port = 16 bits

< 65536

listen(socketfd, 5)

while one client is being serviced by
the server, 5 more clients can wait

TCP server

blocking call

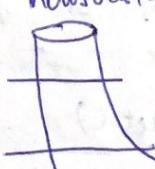
newsocketfd = accept (socketfd, (struct sockaddr *) & __)

(server waits for client's connection)

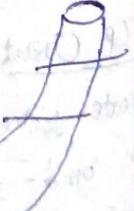
socketfd newsocketfd

client socket

half open socket
2-tuple



4-tuple



```
while(1){
```

```
    newsockfd = accept(...)
```

// communicate with client using newsockfd

```
    close(newsockfd); // newsockfd is gone, but sockfd
```

is still there

```
    {  
        send(sockfd, new buffer, len, 0); }  
        { returns no. of bytes  
        from where? how many bytes. }  
        || recv(sockfd, new buffer, len, 0); }  
        { send or recv (-1  
        where to put? max no. of bytes  
        on error)  
        to receive  
        block if  
        no data in  
        socket buffer
```

TCP guarantees that if I send 100 bytes on one side they will arrive on the other side in the same order but we cannot say that if they arrive altogether.

How to know what I actually received?

Return value of recv call will tell how many bytes we receive.

If recv call comes out to be return value 0 then it means that the other side has closed the connection. This can be used to detect that the other end has sent all the data or not.

TCP client

```
sockfd = socket() - same
```

```
int connect(sockfd, (struct sockaddr *) ...);
```

recv/send
Any application in the world has <IP addr, port>

```
close(sockfd);
```

`servaddr.sin_family = AF_INET;`

`servaddr.sin_port = 20000;`

`servaddr.sin_addr.s_addr = 10.5.18.84;`

(cannot specify this way)

IP address = 32 bit integer

instead there are functions

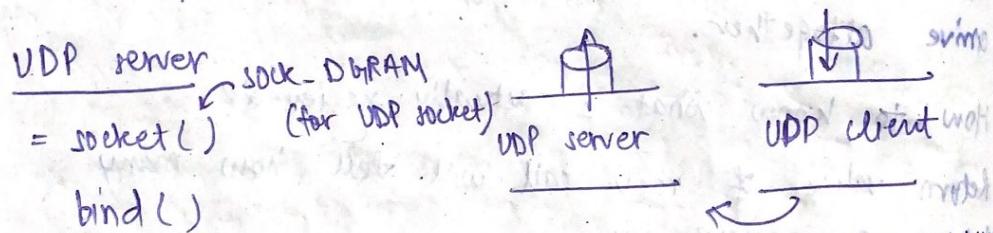
`inet_aton ("10.5.18.84", &serv_addr.sin_addr.s_addr);`

↑
(ASCII to network)

OR

`inetpton (- -)`

Two applications cannot use the same port on the same m/c. Therefore system looks a port for the client and then assigns it (assignment is done by the local m/c). Server always binds & client has an option to bind (may not bind).



replies from
(---, ---, ---, ---, ---, ---)
same client who sent

NO kind of pipes, no reliability

so no kind of accept or

UDP client

`socket()`

`send()` → where I am specifying (where this packet is supposed to go)

We do `sendto (---, ---, ---, ---, ---, ---)`

same

who to send

Here client can send to any server since we are now specifying each time which server to send

In the recvfrom & accept function we have last parameters which specify which client we have received from (in case of UDP) & which client we have connected to (in case of TCP) when ~~the~~ the function comes out.

so, In case of UDP a server first need to make a recvfrom call to get the address of client and then we come to know the client's IP, port. Now, we can use it to send ~~# bad~~ data to client.

All the above calls if they fail the return -1;

perror() fine:

There is a number to every error.

When there is an error, the error no. value is set to some particular value.

perror() ~~is~~ is a func where we can add our own msg then it tells the error after a colon (:)

serv_addr.sin_port = htons(20000);

htons() is used to automatically adapt to the byte orientation used on a particular mc (little Endian OR Big Endian)

There is prescribes standard of network byte order, it says whatever your integer format is, when you put on wire you put it in this order.

htons → host to network short

because port is 16 bits

when the connect call will bind to the accept call

the cli-addr will be filled with the address of the client and the clilen will be filled up with the size of struct sockaddr. So we know which is the client we connected to.

If we give null-pointers to this still it will connect. It just adds up an extra info for the client. 127.0.0.1 is a special address for the local interface. Iterative server - A server which handles client one-by-one. It handles one client in one accept call, then does everything for that client only then it goes to the other accept call again to allow the next client.

Networks:

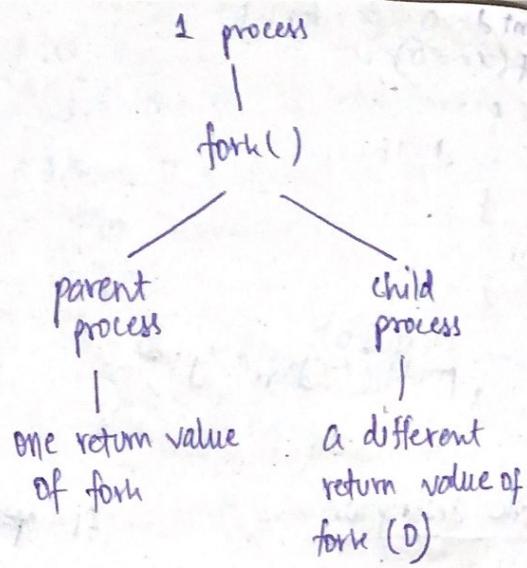
NOW we have a case that we need our server to accept data or send calls from a TCP or a UDP socket whichever comes first. We know that an accept (in TCP) & recvfrom (in UDP) is a blocking call and we don't come out of them until we get a connect (in TCP) & sendto (in UDP). Therefore, we need some other mechanism to block on both TCP & UDP and come out wrt whichever arrives first. This can be done using select/poll.

Also we can have a case of a concurrent server which can process multiple clients at same time. We can handle each request by spawning a new process using fork() call.

int fork()

Program is run in a process. For each program we create a process.

1 process
 |
 function call
 |
 1 process
 1 return value



$x = fork();$ ← 1 process
 if ($x == 0$) ← 2 process
 printf ("This\n");
 else
 printf ("That\n");

Output

This OR That
 That This

int d=10;
 $x = fork();$
 if ($x == 0$)
 {
 gets a copy of d for child process
 }
 else {
 gets a separate copy of d for parent process
 }

Each process gets a copy of all the variables

```

int d=10; x=fork();
if(x==0)
{
    d=15;
}
else
{
    sleep(20);
    printf("%d\n", d);
}

```

File descriptor `f1 = open(...)`

`f1 = open(...)`

file descriptor table		
(keyboard)	0	STDIN
(Display)	1	STDOUT
(Error)	2	STDERR
f1	→	
sockfd	→	
newsockfd	→	

sockfd = socket(...)

Each of open call of whatever you open returns an index to the file descriptor table and the table contains info about the file opened index of last entry

Bit pattern

index 0	0	0	0	-	1	37
↓	↑			↑		
index 1					talking about the "file"	
					at index 37	

FD-SET - bit pattern with 1 bit for every file

FD-SET fd;

Macros do manipulate FD-SETs

FD_SET waitfd;

FD-ZERO (&waitfd); // zero out all the bits

```
FD_SET(&waitfd, sockfd); // Set the bit at index sockfd
FD_CLR(&waitfd, sockfd); // Clears the bit at index sockfd
FD_ISSET(&waitfd, sockfd); // Returns 0 if sockfd is not set
// >0 otherwise
```

```
select (maxfd, &waitfd, 0, 0, &time);  
(int) (FD_SET*) time out
```

timeout is defined as:

```
struct timeval {  
    int tv_sec; // seconds  
    int tv_usec; // microseconds  
};
```

maxfd = max (all sockfd) + 1

it is telling that we only think about bits till
maxfd rest are unimportant (more for optimisation
purpose)

sockfd1, sockfd2

wait on sockfd1 & sockfd2
come out if data in any one or 3 sec have passed

struct timeval t;

FD_SET fd;

FD_ZERO (&fd);

FD_SET (&fd, sockfd1);

FD_SET (&fd, sockfd2);

t.tv_sec = 3;

t.tv_usec = 0;

err = select (max (sockfd1, sockfd2) + 1, &fd, 0, 0, &t);

↓
Blocks until either data in sockfd1, or data in sockfd2,
or 3 sec over

Suppose data comes only in sockfd1, select will wake up.
But we don't know where in which socket we have data.
Here system comes to rescue. The system sets only the bit which has data, the remaining ones will be set to zero.

```
if (FD_ISSET (lfd, sockfd1))  
    /* data in sockfd1 */  
{  
    recv (sockfd1);  
}  
if (FD_ISSET (lfd, sockfd2))  
    /* data in sockfd2 */  
{  
    recv (sockfd2);  
}  
/* timeout */ if (FD_ISSET (lfd, sockfd1) == 0 &  
        FD_ISSET (lfd, sockfd2) == 0)
```

How do you send a series of strings? (String)
(How do you send detect when a string of strings end?)

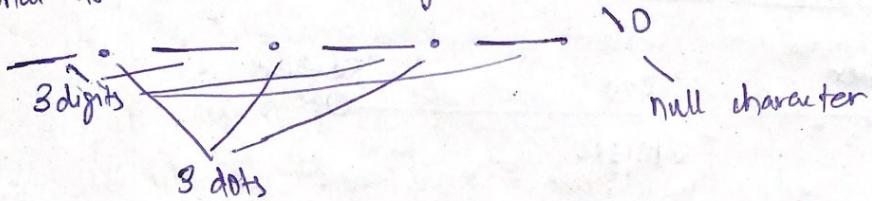
client → server | Client and Listener no. of fields
 | one per n fields for two strings
 ←
 string of strings | Insert with
 "10,14,3,7 \0", "12,5 ---" | T2L-0
 |

We can do this by first sending the number of strings and then the other side can just accept that no. of strings.

→ \0 → \0 → \0\0
| indicator of end of string of strings

Here we have UDP protocol so how can a client know what is the length of the string coming from the server, if the server is sending 300 bytes and client has 100 bytes buffer the 200 bytes will just gets dropped.

So you can't do it in a single PC. But if you are sending IPs (DNS server) then you know what is the max length of an IP address



$$\text{max_size} = 16 \text{ bytes}$$

TCP socket (socketfd1)
block on accept()

UDP socket
block on recvfrom()

```
select (-, -, -)
if (FD_ISSET (&fd-set, socketfd1)
{
    newsocketfd = accept (-, -) // does not block as some
    fork()
    recv()
}
if (FD_ISSET (&fd-set, socketfd2))
```

To see ip addresses of DNS name
nslookup www.google.com

HTTP - Text Based Protocol
Command
Header

fieldname : value

→ GET — 1.1

Connection : close

ACCEPT-LANGUAGE :

ACCEPT-TYPE : application/pdf

Command
Request
Response
Entity

→ PUT — 1.1
≡

How to make send/recv non blocking?

recv(—, —, —, flag)
o

|| MSG_DONTWAIT — for non-blocking

If there is nothing in the buffer, it will fail
return value < 0

ERRNO set to EWOULDBLOCK OR EAGAIN

check ERRNO in if and proceed accordingly.

|| MSG_PEEK

Return data from the beginning of the recv queue.
A subsequent call will return the same data.

|| MSG_WAITALL

This flag indicates that it will block until the full
data request specified in parameters are not received.

How to make all blocking calls non-blocking?

fcntl(sockfd, F_SETVAL, O_NONBLOCK)
F_GETVAL, —

2 bit flags

00 ← Block

10

01 ← Non block

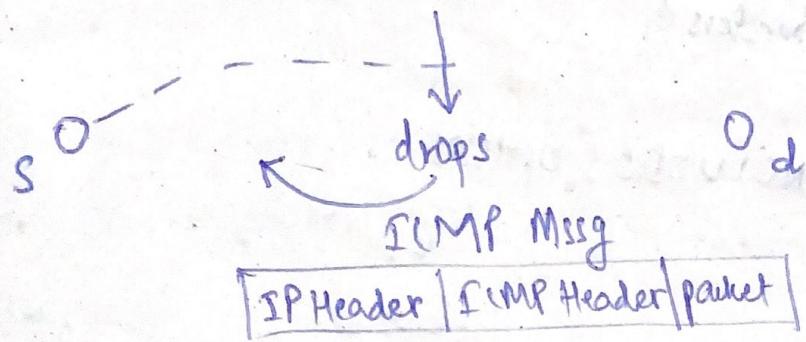
overwritten

This overwrites other flag bits

| in man page F_SETFL & F_GETFL

val = fcntl(sockfd, F_SETVAL, 0);
fcntl(sockfd, F_SETVAL, VAL || O_NONBLOCK);

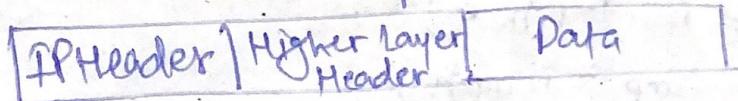
Internet Control Message protocol (ICMP):



destination Unreachable has IP header + 64 bits of IP plet

To know which ~~data~~ plet of IP was dropped

To know about the header & type used in higher layer protocol



Networks Lab:

Gethostbyname ← DNS call

getsockopt

setsockopt

Raw sockets:

socket (___, SOL_SOCKET, IPPROTO_ICMP)

Raw sockets - no ports

sendto/recvfrom

Suppose 5 processes opened 5 raw sockets on IPPROTO-ICMP.

When OS gets an ICMP pkt it gives to all 5 raw sockets.

recvfrom on this will give the entire IP pkt, including IP header & ICMP header

[IP header | ICMP Header | ICMP data]

How to send?

If you don't say anything to the OS, the OS will generate its own IP header.

① allow setting IP header int one=1

setsockopt(sockfd, IPPROTO_IP, IPHDRINCL, &one, sizeof(one))
level

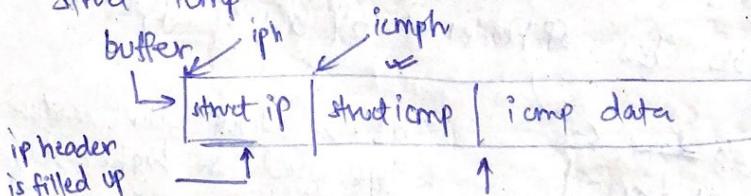
Now how to set these headers?

struct ip

struct ipheader - INCLUDES options

Similarly,

struct icmp



char buff[1000];

iph = (struct ip *) buff;

iph->hlen =

iph->ver =

icmph = buff + sizeof(struct ip) (buff + sizeof(struct ip))

icmph->type =

icmph->code =

char *data = buff + sizeof(struct ip) + sizeof(struct icmp)

sendto(, buff, len, — —)

There are some file ~~for~~ fields which are set automatically like checksum & for some we have options like for identity we can put our own, but if we do not put our own IP will put his own identity.

✓ Lookup alarm() for signal SIGALRM

This is for making the receiver wake up when a message is in the pipe ready to recv.

✓ setsockopt(, SOL_SOCKET, SND_BUFFSIZE, —)
RCV_BUFFSIZE

✓ funtl(, O_ASYNC)

sender

ID, seqno, sending-time

To keep track of lost or dropped plots.